

MC 302– DBMS: Concurrency Control

Goonjan Jain

Dept. of Applied Mathematics

Delhi Technological University

Outline

- Serializability
- Locking –
 - 2PL
 - Variations
- Deadlocks


Formal Properties of Schedules

- Levels of serializability
 - Conflict serializability – all DBMSs support this
 - View serializability – harder but allows more concurrency
- Conflicting operations- Two operations conflict if:
 - They are by different transactions,
 - They are on the same object and at least one of them is a write.


Conflict Serializable Schedules

- Two schedules are ***conflict equivalent*** iff:
 - They involve the same actions of the same transactions, and
 - Every pair of conflicting actions is ordered the same way.
- Schedule S is ***conflict serializable*** if:
 - S is conflict equivalent to some serial schedule.
 - Able to transform S into a serial schedule by swapping consecutive nonconflicting operations of different transactions.

Conflict Serializability Intuition




Conflict Serializability Intuition



Serializability

- **Q:** Are there any faster algorithms to figure this out other than transposing operations?
- **A:** Dependency Graphs
- Dependency Graphs:
 - One node per transaction.
 - Edge from T_i to T_j if:
 - An operation O_i of T_i conflicts with an operation O_j of T_j and
 - O_i appears earlier in the schedule than O_j .
 - Also known as a “precedence graph”
- **Theorem:** A schedule is *conflict serializable* if and only if its dependency graph is acyclic.




Dependency Graphs – Example 1

Schedule

T1	T2
BEGIN	BEGIN
R(A)	
W(A) →	R(A)
	W(A)
	R(B)
	W(B)
	COMMIT
R(B) ←	
W(B)	
COMMIT	

Dependency Graph




Example 2 – Lost update

Schedule


T1	T2
BEGIN R(A) A = A-1 W(A) COMMIT	BEGIN R(A) A = A - 1 W(A) COMMIT

Dependency Graph



Example 3

T1	T2	T3
BEGIN R(A) W(A)		BEGIN R(A) W(A) COMMIT
R(B) W(B) COMMIT	BEGIN R(B) W(B) COMMIT	



Is this equivalent to a serial execution?


- A: Yes (T2, T1, T3)
 - Notice that T3 should go after T2, although it starts before it!
- Need an algorithm for generating serial schedule from an acyclic dependency graph.
 - **Topological Sorting**

Example 4 – Inconsistent Analysis

Schedule

T1	T2
BEGIN	BEGIN
R(A)	
A = A-1	
W(A) ←	R(A) →
	SUM = A
	R(B)
	SUM += B
	ECHO(SUM)
	COMMIT
R(B)	
B = B+1	
W(B) ←	↑
COMMIT	

Dependency Graph




Is it possible to create a schedule similar to this that is “correct” but still not conflict serializable?


View Serializability

- Alternative (weaker) notion of serializability.
- Schedules S_1 and S_2 are ***view equivalent*** if:
 - If T_1 reads initial value of A in S_1 , then T_1 also reads initial value of A in S_2 .
 - If T_1 reads value of A written by T_2 in S_1 , then T_1 also reads value of A written by T_2 in S_2 .
 - If T_1 writes final value of A in S_1 , then T_1 also writes final value of A in S_2 .

View Serializability




View Serializability




Serializability

- **View Serializability** allows (slightly) more schedules than **Conflict Serializability** does.
 - But is difficult to enforce efficiently.
- In practice, **Conflict Serializability** is what gets used, because it can be enforced efficiently.




Locking Based Concurrency Control

- Lost update problem – without locks



Executing with locks

- With locks – lock manager grants/denies lock requests



Executing with locks

- **Q:** If a transaction only needs to read 'A', should it still get a lock?
- **A:** Yes, but you can get a shared lock.


Lock Types

- Basic Lock Types:
 - S-Lock: Shared Locks (Reads)
 - X-Lock: Exclusive Locks (writes)

Compatibility Matrix			
T1 has	T2 wants	Shared	Exclusive
Shared	Shared	Y	N
Exclusive	Exclusive	N	N

Executing with locks

- Transactions request locks (or upgrades)
 - Lock manager grants or blocks requests
 - Transactions release locks
 - Lock manager updates lock-table
- *But this is not enough...*
- *Inconsistent Analysis*




Concurrency Control

- We need to use a **well-defined protocol** that ensures that transactions execute correctly.
- Two categories:
 - Two-Phase Locking (2PL)
 - Timestamp Ordering (T/O) – discuss in future classes

Two-Phase Locking

- **Phase 1: Growing**
 - Each transaction requests the locks that it needs from the DBMS's lock manager.
 - The lock manager grants/denies lock requests.
- **Phase 2: Shrinking**
 - The transaction is allowed to only release locks that it previously acquired.
 - It cannot acquire new locks.
- The transaction is not allowed to acquire/upgrade locks after the growing phase finishes.

Executing with 2PL



Lock Management

- Lock and unlock requests handled by the DBMS's *lock manager* (LM).
- LM contains an entry for each currently held lock:
 - Pointer to a list of transactions holding the lock.
 - The type of lock held (shared or exclusive).
 - Pointer to queue of lock requests.
- When lock request arrives see if any other transaction holds a conflicting lock.
 - If not, create an entry and grant the lock
 - Else, put the requestor on the wait queue
- All lock operations must be atomic.
- Lock upgrade: The transaction that holds a shared lock upgrade to hold an exclusive lock.

Two-Phase Locking

- 2PL –
 - sufficient to guarantee conflict serializability (i.e., precedence graph is acyclic),
 - But, subject to *cascading aborts*.

Schedule	
T1	T2
BEGIN	BEGIN
X-LOCK(A)	
X-LOCK(B)	
R(A)	
W(A)	
UNLOCK(A)	
	X-LOCK(A)
R(B)	R(A)
W(B)	W(A)
ABORT	:

This is a permissible schedule in 2PL, but we have to abort T2 too.

This is all wasted work!


2PL Observations

- There are schedules that are serializable but would not be allowed by 2PL.
- Locking limits concurrency.
- May lead to deadlocks.
- May still have “dirty reads”
 - Solution: **Strict 2PL**


Strict Two-Phase Locking

- The transaction is not allowed to acquire/upgrade locks after the growing phase finishes.
- Allows only conflict serializable schedules, but it is actually stronger than needed.
- A schedule is ***strict*** if a value written by a transaction is not read or overwritten by other transactions until that transaction finishes.
- Advantages:
 - Recoverable.
 - Do not require cascading aborts.
 - Aborted transactions can be undone by just restoring original values of modified tuples.


Non 2PL Example



2PL Example




Strict 2PL Example




Strict Two-Phase Locking

- Transactions hold all of their locks until commit.
- Good:
 - Avoids “dirty reads” etc.
- Bad:
 - Limits concurrency even more
 - And still may lead to deadlocks

Schedules






Deadlocks

- **Deadlock:** Cycle of transactions waiting for locks to be released by each other.
- Two ways of dealing with deadlocks:
 - Deadlock prevention
 - Deadlock detection
- Many systems just punt and use timeouts
 - What are the dangers with this approach?

Deadlock Detection

- The DBMS creates a ***waits-for*** graph:
 - Nodes are transactions
 - Edge from T_i to T_j if T_i is waiting for T_j to release a lock
- The system periodically check for cycles in ***waits-for*** graph.

T1	T2	T3
BEGIN S-LOCK(A) S-LOCK(D)	BEGIN X-LOCK(B)	BEGIN S-LOCK(C)
S-LOCK(B)	X-LOCK(C)	X-LOCK(A)



Deadlock Detection

- What do we do when we find a deadlock?
- Select a “victim” and rollback it back to break the deadlock.
- Which one do we choose?
- It depends...
 - By age (lowest timestamp)
 - By progress (least/most queries executed)
 - By the # of items already locked
 - By the # of transactions that we have to rollback with it
- We also should consider the # of times a transaction has been restarted in the past.
- How far do we rollback?
- It depends...
 - Completely
 - Minimally (i.e., just enough to release locks)


Deadlock Prevention

- When a transaction tries to acquire a lock that is held by another transaction, kill one of them to prevent a deadlock.
- Assign priorities based on timestamps:
 - Older → higher priority (e.g., $T_1 > T_2$)
- Two different prevention policies:
 - **Wait-Die:** If T_1 has higher priority, T_1 waits for T_2 ; otherwise T_1 aborts (“old wait for young”)
 - **Wound-Wait:** If T_1 has higher priority, T_2 aborts; otherwise T_1 waits (“young wait for old”)

Deadlock Prevention


Wait-Die

T_i wants T_j has



Wound-Wait

T_i wants T_j has



Deadlock Prevention

T1	T2
BEGIN X-LOCK(A)	BEGIN X-LOCK(A)



Wait-Die

T1 Waits

Wound-Wait

T2 aborted

T1	T2
BEGIN X-LOCK(A)	BEGIN X-LOCK(A)



Wait-Die

T2 Aborted

Wound-Wait

T1 waits

Deadlock Prevention

- Why do these schemes guarantee no deadlocks?
- Only one “type” of direction allowed.
- When a transaction restarts, what is its (new) priority?
- Its original timestamp. Why?

Performance Problems


- Executing more transactions can increase the throughput.
- But there is a tipping point where adding more transactions actually makes performance worse.

Quiz

- is there a serial schedule (= interleaving) that is not serializable?
- is there a serializable schedule that is not serial?
- can 2PL produce a non-serializable schedule? (assume no deadlocks)
- is there a serializable schedule that can not be produced by 2PL?

Lock Granularities

- When we say that a transaction acquires a “lock”, what does that actually mean?
 - On a field? Record? Page? Table?
- Ideally, each transaction should obtain fewest number of locks that is needed...



Intention Locks

- Intention locks allow a higher level node to be locked in **S** or **X** mode without having to check all descendent nodes.
- If a node is in an intention mode, then explicit locking is being done at a lower level in the tree.

Intention Locks


- Useful in practice as each transaction only needs a few locks.
- Intention locks help improve concurrency:
- **Intention-Shared (IS)**: Indicates explicit locking at a lower level with shared locks.
- **Intention-Exclusive (IX)**: Indicates locking at lower level with exclusive or shared locks.
- **Shared+Intention-Exclusive (SIX)**: The subtree rooted by that node is locked explicitly in shared mode and explicit locking is being done at a lower level with exclusive-mode locks.

Compatibility Matrix


		T2 Wants				
		IS	IX	S	SIX	X
T1 Holds	IS	✓	✓	✓	✓	✗
	IX	✓	✓	✗	✗	✗
	S	✓	✗	✓	✗	✗
	SIX	✓	✗	✗	✗	✗
	X	✗	✗	✗	✗	✗

Multiple Granularity Protocol

- Each transaction obtains appropriate lock at highest level of the database hierarchy.
- To get **S** or **IS** lock on a node, the txn must hold at least **IS** on parent node.
 - What if transaction holds **SIX** on parent?
S on parent?
- To get **X**, **IX**, or **SIX** on a node, must hold at least **IX** on parent node.




Example – Two Level Hierarchy



Example

- Assume three transactions execute at same time:
- **T1:** Scan R and update a few tuples.
- **T2:** Scan a portion of tuples in R.
- **T3:** Scan all tuples in R.



Example

- **T1:** Get an **SIX** lock on **R**, then get **X** lock on tuples that are updated.
- **T2:** Get an **IS** lock on **R**, and repeatedly get an **S** lock on tuples of **R**.
- **T3:** Two choices:
 - T3 gets an **S** lock on **R**.
 - OR, T3 could behave like T2; can use **lock escalation** to decide which.

